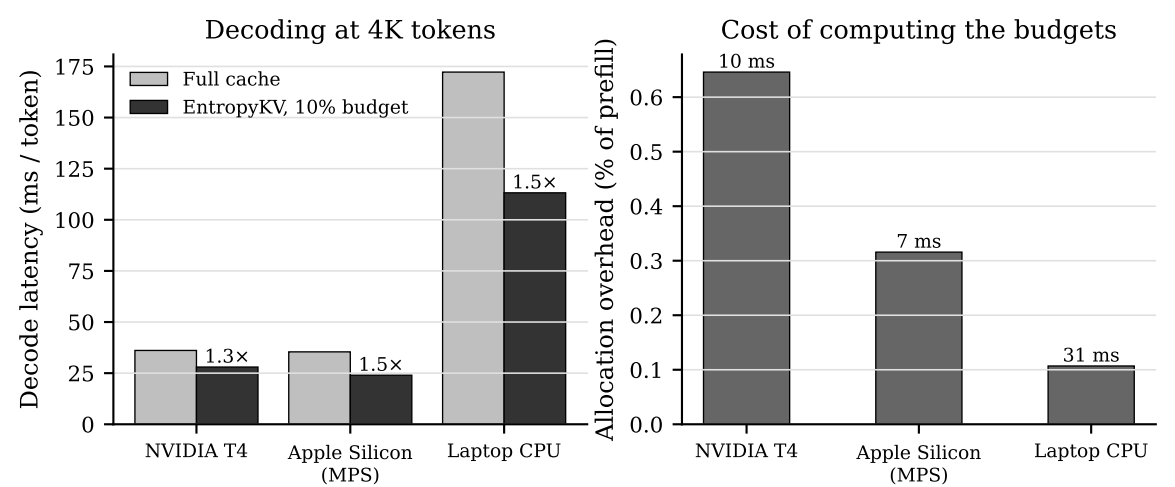
Decoding at 4K tokens
Decode latency (ms / token)
175
150
125
100
75
50
25
0
Full cache
EntropyKV, 10% budget
1.3×
1.5×
1.5×
NVIDIA T4
Apple Silicon (MPS)
Laptop CPU
Cost of computing the budgets
Allocation overhead (% of prefill)
0.6
0.5
0.4
0.3
0.2
0.1
0.0
10 ms
7 ms
31 ms
NVIDIA T4
Apple Silicon (MPS)
Laptop CPU